

# hw8

2024-03-29

## 2a

```
library(dplyr)

##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(MASS)

##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##   select

Y <- read.table("http://www2.stat.duke.edu/~pdh10/FCBS/Exercises/azdiabetes.dat", header=TRUE)
diabetics <- Y %>%
  filter(diabetes=="Yes")
Y_d <- diabetics[, 1:7]
nondiabetics <- Y %>%
  filter(diabetes=="No")
Y_n <- nondiabetics[, 1:7]

mu0_d <- colMeans(diabetics[, 1:7])
samp_cov_d <- cov(diabetics[, 1:7])
lambda0_d <- cov(diabetics[, 1:7])
S0_d <- samp_cov_d
nu0_d <- 9

mu0_n <- colMeans(nondiabetics[, 1:7])
samp_cov_n <- cov(nondiabetics[, 1:7])
lambda0_n <- cov(nondiabetics[, 1:7])
S0_n <- samp_cov_n
nu0_n <- 9

### Simulate multivariate normal vector
rmvnorm<-function(n,mu,Sigma)
{
  p<-length(mu)
```

```

res<-matrix(0,nrow=n,ncol=p)
if(n>0 & p>0)
{
  E<-matrix(rnorm(n*p),n,p)
  res<-t( t(E%%chol(Sigma)) +c(mu))
}
res
}

### Simulate from the Wishart distribution
rwish<-function(n,nu0,S0)
{
  sS0 <- chol(S0)
  S<-array( dim=c( dim(S0),n ) )
  for(i in 1:n)
  {
    Z <- matrix(rnorm(nu0 * dim(S0)[1]), nu0, dim(S0)[1]) %%% sS0
    S[,i]<- t(Z)%%Z
  }
  S[,1:n]
}

### Gibbs sampler

Sigma_d <- samp_cov_d
n_d<-dim(Y_d)[1]
S0_d <- samp_cov_d
THETA_d <- NULL
SIGMA_d <- NULL

Sigma_n <- samp_cov_n
n_n<-dim(Y_n)[1]
S0_n <- samp_cov_n
THETA_n <- NULL
SIGMA_n <- NULL

S <- 10000

for (s in 1:S){
  ###update theta_d
  Ln_d<-solve( solve(lambda0_d) + n_d*solve(Sigma_d) )
  mun_d<-Ln_d%%( solve(lambda0_d)%%mu0_d + n_d*solve(Sigma_d)%%mu0_d )
  theta_d<-rmvnorm(1,mun_d,Ln_d)
  ###

  ###update Sigma_d
  Sn_d<- S0_d + ( t(Y_d)-c(theta_d) )%%t( t(Y_d)-c(theta_d) )
  Sigma_d<-solve( rwish(1, nu0_d+n_d, solve(Sn_d)) )
  ###

  ### save results
  THETA_d<-rbind(THETA_d,theta_d) ; SIGMA_d<-rbind(SIGMA_d,c(Sigma_d))
  ###
}

```

```

    ###update theta_n
Ln_n<-solve( solve(lambda0_n) + n_n*solve(Sigma_n) )
mun_n<-Ln_n%*%( solve(lambda0_n)%*%mu0_n + n_n*solve(Sigma_n)%*%mu0_n )
theta_n<-rmvnorm(1,mun_n,Ln_n)
###

    ###update Sigma_n
Sn_n<- S0_n + ( t(Y_n)-c(theta_n) )%*%t( t(Y_n)-c(theta_n) )
Sigma_n<-solve( rwish(1, nu0_d+n_n, solve(Sn_n)) )
###

    ### save results
THETA_n<-rbind(THETA_n,theta_n) ; SIGMA_n<-rbind(SIGMA_n,c(Sigma_n))
###

if (s %% 100 == 0){
  print(s)
}

}

```

```

## [1] 100
## [1] 200
## [1] 300
## [1] 400
## [1] 500
## [1] 600
## [1] 700
## [1] 800
## [1] 900
## [1] 1000
## [1] 1100
## [1] 1200
## [1] 1300
## [1] 1400
## [1] 1500
## [1] 1600
## [1] 1700
## [1] 1800
## [1] 1900
## [1] 2000
## [1] 2100
## [1] 2200
## [1] 2300
## [1] 2400
## [1] 2500
## [1] 2600
## [1] 2700
## [1] 2800
## [1] 2900

```

```
## [1] 3000
## [1] 3100
## [1] 3200
## [1] 3300
## [1] 3400
## [1] 3500
## [1] 3600
## [1] 3700
## [1] 3800
## [1] 3900
## [1] 4000
## [1] 4100
## [1] 4200
## [1] 4300
## [1] 4400
## [1] 4500
## [1] 4600
## [1] 4700
## [1] 4800
## [1] 4900
## [1] 5000
## [1] 5100
## [1] 5200
## [1] 5300
## [1] 5400
## [1] 5500
## [1] 5600
## [1] 5700
## [1] 5800
## [1] 5900
## [1] 6000
## [1] 6100
## [1] 6200
## [1] 6300
## [1] 6400
## [1] 6500
## [1] 6600
## [1] 6700
## [1] 6800
## [1] 6900
## [1] 7000
## [1] 7100
## [1] 7200
## [1] 7300
## [1] 7400
## [1] 7500
## [1] 7600
## [1] 7700
## [1] 7800
## [1] 7900
## [1] 8000
## [1] 8100
## [1] 8200
## [1] 8300
```

```
## [1] 8400
## [1] 8500
## [1] 8600
## [1] 8700
## [1] 8800
## [1] 8900
## [1] 9000
## [1] 9100
## [1] 9200
## [1] 9300
## [1] 9400
## [1] 9500
## [1] 9600
## [1] 9700
## [1] 9800
## [1] 9900
## [1] 10000
```

```
colMeans(THETA_n) - colMeans(THETA_d)
```

```
##      npreg      glu      bp      skin      bmi      ped
## -1.7756579 -33.0601339 -4.8013461 -5.6817419 -4.3842703 -0.1699392
##      age
## -7.2096250
```

```
apply(THETA_n, 2, var) - apply(THETA_d, 2, var)
```

```
##      npreg      glu      bp      skin      bmi
## -0.0653479875 -3.9133928337 -0.4587850935 -0.3142579553 -0.1214909829
##      ped      age
## -0.0006571537 -0.3856434809
```

The glucose of diabetics seems to be, on average, much higher than that of non-diabetics. Additionally, the variance of glucose is much greater for diabetics than it is for non-diabetics.

```
colMeans(THETA_d > THETA_n)
```

```
## npreg  glu  bp  skin  bmi  ped  age
##      1   1   1   1   1   1   1
```

$Pr(\theta_{d,j} > \theta_{n,j}|Y) = 1$  for all  $j \in \{1, 2, 3, 4, 5, 6, 7\}$ .


## 2b

```
colMeans(SIGMA_n) - colMeans(SIGMA_d)
```

```
## [1] -7.60154621 14.38438901 0.51952251 7.77626691 4.65362092
## [6] 0.05342474 -5.23881308 14.38438901 -386.84365430 22.92976838
## [11] 1.36213178 15.15711759 0.41321728 7.73697657 0.51952251
## [16] 22.92976838 -15.05849430 16.29929566 5.05912658 0.04328221
## [21] 3.17189992 7.77626691 1.36213178 16.29929566 -6.54805383
## [26] 8.81520427 -0.48356095 24.91181706 4.65362092 15.15711759
## [31] 5.05912658 8.81520427 -0.76124061 -0.29514813 18.24530299
## [36] 0.05342474 0.41321728 0.04328221 -0.48356095 -0.29514813
## [41] -0.06985385 0.21283050 -5.23881308 7.73697657 3.17189992
## [46] 24.91181706 18.24530299 0.21283050 -19.30694013
```

The 2,2 entry corresponds to glucose, so this supports our observation that the variability in glucose is especially high in diabetics.

```
plot(colMeans(SIGMA_n), colMeans(SIGMA_d),
     xlab = "Posterior Mean of Covariance Matrix 1",
     ylab = "Posterior Mean of Covariance Matrix 2",
     main = "Posterior Means of Covariance Matrices",
     pch = 19, col = "blue")
line <- lm(colMeans(SIGMA_d)~colMeans(SIGMA_n))
abline(line)
```



The entries of the two covariance matrices are positively correlated. This visualization makes it apparent that the variance of glucose is high for both groups when compared to that of the other covariates.

## 4b

```
Y<-dget(url("http://www2.stat.duke.edu/~pdh10/FCBS/Inline/Y.pima.miss"))
```

```
Y=Y%>%
  filter(!is.na(glu))%>%
  filter(!is.na(bp))%>%
  filter(!is.na(skin))%>%
  filter(!is.na(bmi))
colMeans(Y)
```

```
##      glu      bp      skin      bmi
## 121.85039  70.61417  28.83465  31.67874
```

```

###
Y <- readRDS("hw8train.rds")

### prior parameters
n<-dim(Y)[1] ; p<-dim(Y)[2]
mu0<-c(rep(0,14))
sd0<-(mu0/2)
L0<-matrix(0,p,p) ; diag(L0)<-1 # ; L0<-L0*outer(sd0,sd0)
nu0<-p+2 ; S0<-L0
###

### starting values
Sigma<-S0
Y.full<-Y
O<-1*(!is.na(Y))
for(j in 1:p)
{
  Y.full[is.na(Y.full[,j]),j]<-mean(Y.full[,j],na.rm=TRUE)
}
###

### Gibbs sampler
THETA<-SIGMA<-Y.MISS<-NULL
set.seed(1)
S<-1000
for(s in 1:S)
{

  ###update theta
  ybar<-apply(Y.full,2,mean)
  Ln<-solve( solve(L0) + n*solve(Sigma) )
  mun<-Ln*%( solve(L0)*%mu0 + n*solve(Sigma)*%ybar )
  theta<-rmvnorm(1,mun,Ln)
  ###

  ###update Sigma
  Sn<- S0 + ( t(Y.full)-c(theta) )%*%t( t(Y.full)-c(theta) )
  Sigma<-solve( rwish(1, nu0+n, solve(Sn)) )
  ###

  ###update missing data
  for(i in 60:n)
  {
    b <- ( O[i,]==0 )
    #print(b)
    a <- ( O[i,]==1 )
    #print(a)
    iSa<- solve(Sigma[a,a])
    beta.j <- Sigma[b,a]*%iSa
    Sigma.j <- Sigma[b,b] - Sigma[b,a]*%iSa*%Sigma[a,b]
    #print(dim(beta.j))
    #print(dim(t(Y.full[i,a])))
    #print(theta)
    #print(dim(theta))
  }
}

```

```

    #print(theta[a])
    #print(dim(theta[a]))
    #print(matrix(theta[a]))
    theta.j<- matrix(theta[b]) + beta.j%*%matrix((t(Y.full[i,a])-theta[a]))
    Y.full[i,b] <- rmvnorm(1,theta.j,Sigma.j )
  }

  ### save results
  THETA<-rbind(THETA,theta) ; SIGMA<-rbind(SIGMA,c(Sigma))
  Y.MISS<-rbind(Y.MISS, Y.full[0==0] )
  ###
}

```

## 4c

```

## compare to test dataset
Y.true<-readRDS("hw8test.rds")





V<-matrix(1:p,nrow=n,ncol=p,byrow=TRUE)

v.miss<-V[0==0]
y.pred<-apply(Y.MISS,2,mean)
y.true<-Y.true#[0==0]
par(mfrow=c(2,2),mar=c(3,3,1,1),mgp=c(1.75,.75,0))
for(j in 8:p){
  #print(y.pred[v.miss==j])
  #print(y.true[v.miss==j])
  plot(y.true[v.miss==j], y.pred[v.miss==j],
       xlab=paste("true", colnames(Y.true)[j]),
       ylab=paste("predictied", colnames(Y.true)[j]),pch=16,
       xlim=range(c(y.pred[v.miss==j],y.true[v.miss==j])),
       ylim=range(c(y.pred[v.miss==j],y.true[v.miss==j])))
  abline(0,1)
  cat(j, mean( (y.true[v.miss==j]- y.pred[v.miss==j])^2),
      mean( (y.true[v.miss==j]- mean(Y[,j],na.rm=TRUE))^2),"\n")
}




## 8 1.25844 2.521915
## 9 1.476817 2.600428
## 10 0.7580454 1.963907

```





```
## 11 0.9825413 2.7985
## 12 1.207481 2.759375
## 13 0.780527 1.715279
## 14 1.394759 2.389352
```



The error of our predicted values is less than the error of  $\hat{\theta}_B$ .